

In this notebook, You will do amazon review classification with BERT.[Download data from [this link](#)]

It contains 5 parts as below. Detailed instructions are given in the each cell. please read every comment we have written.

1. Preprocessing
2. Creating a BERT model from the Tensorflow HUB.
3. Tokenization
4. getting the pretrained embedding Vector for a given review from the BERT.
5. Using the embedding data apply NN and classify the reviews.
6. Creating a Data pipeline for BERT Model.

## instructions:

1. Don't change any Grader Functions. Don't manipulate any Grader functions. If you manipulate any, it will be considered as plagiarised.
2. Please read the instructions on the code cells and markdown cells. We will explain what to write.
3. please return outputs in the same format what we asked. Eg. Don't return List if we are asking for a numpy array.
4. Please read the external links that we are given so that you will learn the concept behind the code that you are writing.
5. We are giving instructions at each section if necessary, please follow them.

Every Grader function has to return True.

In [ ]:

```
!pip install ipython-autotime
%load_ext autotime
```

```
Collecting ipython-autotime
  Downloading
https://files.pythonhosted.org/packages/e6/f9/0626bbdb322e3a078d968e87e3b01341e7890544de891d0cb61360e6/ipython-autotime-0.1.tar.bz2
Building wheels for collected packages: ipython-autotime
  Building wheel for ipython-autotime (setup.py) ... done
  Created wheel for ipython-autotime: filename=ipython-autotime-0.1-cp36-none-any.whl size=1831 sha256=9aad8a1eef47becf5e1c229e797f501a1306a756b710f0ee4c7b9aed4e97e
  Stored in directory:
/root/.cache/pip/wheels/d2/df/81/2db1e54bc91002cec40334629bc39cfa86dff540b304ebcd6e
Successfully built ipython-autotime
Installing collected packages: ipython-autotime
Successfully installed ipython-autotime-0.1
```

In [ ]:

```
#all imports
import numpy as np
import pandas as pd
import tensorflow as tf
import tensorflow_hub as hub
from tensorflow.keras.models import Model
```

time: 1.82 s

In [ ]:

```
tf.test.gpu_device_name()
```

```
Out[ ]:
'/device:GPU:0'
```

time: 5.34 s

## Grader function 1

```
In [ ]:
```

```
def grader_tf_version():
    assert((tf.__version__)>'2')
    return True
grader_tf_version()
```

```
Out[ ]:
```

True

time: 7.92 ms

## Part-1: Preprocessing



```
In [ ]:
```

```
!wget --header="Host: storage.googleapis.com" --header="User-Agent: Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/84.0.4147.135 Safari/537.36" --header="Accept: text/html,application/xhtml+xml,application/xml;q=0.9,image/webp,image/apng,*/*;q=0.8,application/signed-exchange;v=b3;q=0.9" --header="Accept-Language: en-IN,en-GB;q=0.9,en-US;q=0.8,en;q=0.7" --header="Referer: https://www.kaggle.com/" "https://storage.googleapis.com/kaggle-data-sets/18%2F2157%2Fbundle%2Farchive.zip?GoogleAccessId=gcp-kaggle-com@kaggle-161607.iam.gserviceaccount.com&Expires=1598752807&Signature=QzbzV8Q1GXtivMKx7XiYaOKDMA sidky5HnYMMc(QbmSTTU1AZh95esx1IBDuTTIzmQ3yHbmwUejOzeyjC8cIJR17CXWbk8TS2glaqOqqq4i0vOK9plymr9X%2Bdpu%2BFwuhwX2EBIvlnN4bYmMkt7M4q60GS5QhswIuJHW06cw8gEj04Aqcxu%2FGaYGIyoN3nhlUiiAOuF2nQ9QYY9%2F3W0LqftuHO1o3iKu416vxNB06WN2XWtcAnl08GpXeRNZ3sSRMAfbDpj6nKR8jITyv8HfIm%2B01f%2FO23NQ15LxtC1CwGGB5jpejSn%2FRIpjNLP7CJyoMMKpHvA%3D%3D" -c -O '18_2157_bundle_archive.zip'
```

```
--2020-08-27 02:01:24-- https://storage.googleapis.com/kaggle-data-sets/18%2F2157%2Fbundle%2Farchive.zip?GoogleAccessId=gcp-kaggle-com@kaggle-161607.iam.gserviceaccount.com&Expires=1598752807&Signature=QzbzV8Q1GXtivMKx7XiYaOKDMA sidky5HnYMMc(QbmSTTU1AZh95esx1IBDuTTIzmQ3yHbmwUejOzeyjC8cIJR17CXWbk8TS2glaqOqqq4i0vOK9plymr9X%2Bdpu%2BFwuhwX2EBIvlnN4bYmMkt7M4q60GS5QhswIuJHW06cw8gEj04Aqcxu%2FGaYGIyoN3nhlUiiAOuF2nQ9QYY9%2F3W0LqftuHO1o3iKu416vxNB06WN2XWtcAnl08GpXeRNZ3sSRMAfbDpj6nKR8jITyv8HfIm%2B01f%2FO23NQ15LxtC1CwGGB5jpejSn%2FRIpjNLP7CJyoMMKpHvA%3D%3D
Resolving storage.googleapis.com (storage.googleapis.com)... 64.233.189.128, 108.177.97.128, 2404:6800:4008:c04::80, ...
Connecting to storage.googleapis.com (storage.googleapis.com)|64.233.189.128|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 253873708 (242M) [application/zip]
Saving to: '18_2157_bundle_archive.zip'
```

18\_2157\_bundle\_arch 100%[=====>] 242.11M 20.9MB/s in 12s

2020-08-27 02:01:37 (20.9 MB/s) - '18\_2157\_bundle\_archive.zip' saved [253873708/253873708]

time: 13 s

```
In [ ]:
```

```
!unzip 18_2157_bundle_archive.zip
```

```
Archive: 18_2157_bundle_archive.zip
  inflating: Reviews.csv
  inflating: database.sqlite
  inflating: hashes.txt
time: 7.2 s
```

In [ ]:

```
#Read the dataset - Amazon fine food reviews
reviews = pd.read_csv("/content/Reviews.csv")
#check the info of the dataset
reviews.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 568454 entries, 0 to 568453
Data columns (total 10 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   Id                    568454 non-null  int64
 1   ProductId            568454 non-null  object
 2   UserId               568454 non-null  object
 3   ProfileName          568438 non-null  object
 4   HelpfulnessNumerator  568454 non-null  int64
 5   HelpfulnessDenominator 568454 non-null  int64
 6   Score                568454 non-null  int64
 7   Time                 568454 non-null  int64
 8   Summary              568427 non-null  object
 9   Text                 568454 non-null  object
dtypes: int64(5), object(5)
memory usage: 43.4+ MB
time: 3.48 s
```

In [ ]:

```
#get only 2 columns - Text, Score
#drop the NAN values

reviews=reviews[['Text','Score']]
```

time: 75.8 ms

In [ ]:

```
reviews.dropna()
reviews.head(2)
```

Out[ ]:

	Text	Score
0	I have bought several of the Vitality canned d...	5
1	Product arrived labeled as Jumbo Salted Peanut...	1

time: 156 ms

In [ ]:

```
#if score> 3, set score = 1
#if score<=2, set score = 0
#if score == 3, remove the rows.

reviews.loc[reviews['Score'] < 3, 'Score'] = 0
reviews.loc[reviews['Score'] > 3, 'Score'] = 1
reviews = reviews.drop(reviews[reviews.Score == 3].index)
```

time: 78 ms

## Grader function 2

In [ ]:

```
def grader_reviews():
```

```
temp_shape = (reviews.shape == (525814, 2)) and (reviews.Score.value_counts()[1]==443777)
assert(temp_shape == True)
return True
grader_reviews()
```

Out[ ]:

True

time: 20.2 ms

In [ ]:

```
def get_wordlen(x):
    return len(x.split())
reviews['len'] = reviews.Text.apply(get_wordlen)
reviews = reviews[reviews.len<50]
reviews = reviews.sample(n=100000, random_state=30)
```

time: 2.42 s

In [ ]:

```
#remove HTML from the Text column and save in the Text column only
import re
for i in reviews['Text']:
    i=re.sub("[<\([\].*?(\)\>]", "", i)
```

time: 268 ms

In [ ]:

```
#print head 5
reviews.head(5)
```

Out[ ]:

	Text	Score	len
64117	The tea was of great quality and it tasted lik...	1	30
418112	My cat loves this. The pellets are nice and s...	1	31
357829	Great product. Does not completely get rid of ...	1	41
175872	This gum is my favorite! I would advise every...	1	27
178716	I also found out about this product because of...	1	22

time: 22.7 ms

In [ ]:

```
#split the data into train and test data(20%) with Stratify sampling, random state 33,
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(reviews["Text"], reviews["Score"],stratify=reviews["Score"],train_size=0.80)
```

time: 510 ms

In [ ]:

```
#plot bar graphs of y_train and y_test
import matplotlib.pyplot as plt
import seaborn as sns
y_train_plot=y_train.to_frame()
y_test_plot=y_test.to_frame()
```



```

sns.countplot(x='Score', data = y_train_plot)
plt.title("Train")
plt.show()

sns.countplot(x='Score', data = y_test_plot)
plt.title("Test")
plt.show()

```

/usr/local/lib/python3.6/dist-packages/statsmodels/tools/\_testing.py:19: FutureWarning: pandas.util.testing is deprecated. Use the functions in the public API at pandas.testing instead.  
import pandas.util.testing as tm



time: 366 ms

In [ ]:

```

#saving to disk. if we need, we can load preprocessed data directly.
reviews.to_csv('preprocessed.csv', index=False)

```

time: 1.65 s

## Part-2: Creating BERT Model

If you want to know more about BERT, You can watch live sessions on Transformers and BERT. we will strongly recommend you to read [Transformers](#), [BERT Paper](#) and, [This blog](#).

For this assignment, we are using [BERT uncased Base model](#).

It uses L=12 hidden layers (i.e., Transformer blocks), a hidden size of H=768, and A=12 attention heads.

In [ ]:

```

## Loading the Pretrained Model from tensorflow HUB
tf.keras.backend.clear_session()

# maximum length of a seq in the data we have, for now i am making it as 55. You can change this
max_seq_length = 55

#BERT takes 3 inputs

#this is input words. Sequence of words represented as integers
input_word_ids = tf.keras.layers.Input(shape=(max_seq_length,), dtype=tf.int32, name="input_word_ids")

#mask vector if you are padding anything
input_mask = tf.keras.layers.Input(shape=(max_seq_length,), dtype=tf.int32, name="input_mask")

#segment vectors. If you are giving only one sentence for the classification, total seg vector is 0.
#If you are giving two sentences with [sep] token separated, first seq segment vectors are zeros and
#second seq segment vector are 1's
segment_ids = tf.keras.layers.Input(shape=(max_seq_length,), dtype=tf.int32, name="segment_ids")

#bert layer
bert_layer = hub.KerasLayer("https://tfhub.dev/tensorflow/bert_en_uncased_L-12_H-768_A-12/1",
trainable=False)
pooled_output, sequence_output = bert_layer([input_word_ids, input_mask, segment_ids])

#Bert model
#We are using only pooled output not sequence out.
#If you want to know about those, please read https://www.kaggle.com/questions-and-answers/86510
bert_model = Model(inputs=[input_word_ids, input_mask, segment_ids], outputs=pooled_output)

```

time: 18.8 s

In [ ]:

```
bert_model.summary()
```

Model: "functional\_1"

Layer (type)	Output Shape	Param #	Connected to
input_word_ids (InputLayer)	[(None, 55)]	0	
input_mask (InputLayer)	[(None, 55)]	0	
segment_ids (InputLayer)	[(None, 55)]	0	
keras_layer (KerasLayer)	[(None, 768), (None, 109482241		input_word_ids[0][0] input_mask[0][0] segment_ids[0][0]
=====			
Total params: 109,482,241			
Trainable params: 0			
Non-trainable params: 109,482,241			

time: 20.9 ms

In [ ]:

```
bert_model.output
```

Out[ ]:

```
<tf.Tensor 'keras_layer/StatefulPartitionedCall:0' shape=(None, 768) dtype=float32>
```

time: 6.03 ms

In [ ]:

```
#getting Vocab file
vocab_file = bert_layer.resolved_object.vocab_file.asset_path.numpy()
do_lower_case = bert_layer.resolved_object.do_lower_case.numpy()
```

time: 9.16 ms

In [ ]:

```
# While importing tokenization , it throws an error so, we have install the package sentencepiece
# ModuleNotFoundError: No module named 'sentencepiece'
!pip install sentencepiece
```

Collecting sentencepiece

Downloading

https://files.pythonhosted.org/packages/d4/a4/d0a884c4300004a78cca907a6ff9a5e9fe4f090f5d95ab341c53dc58/sentencepiece-0.1.91-cp36-cp36m-manylinux1\_x86\_64.whl (1.1MB)  
|██| 1.1MB 2.8MB/s

Installing collected packages: sentencepiece

Successfully installed sentencepiece-0.1.91

time: 3.38 s

In [ ]:

```
!wget --header="Host: doc-14-5k-docs.googleusercontent.com" --header="User-Agent: Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/84.0.4147.135 Safari/537.36" --header="Accept: text/html,application/xhtml+xml,application/xml;q=0.9,image/webp,image/apng,*/*;q=0.8,application/signed-exchange;v=b3;q=0.9" --header="Accept-Language: en-IN,en-GB;q=0.9,en-US;q=0.8,en;q=0.7" --header="Referer: https://drive.google.com/drive/folders/1RBD3mM5Ea5aH6HS6d32glqDMOJA0ciLm" --header="Cookie: AUTH_gnb78hmdiks9t0b8kec09hpa7nncs5e_nonce=79hmk03b4fg0e; _ga=GA1.2.1804417035.1594643089; NID=204=WE4tnDNREeWA-i7VvKayXZRT0nZVp1JeYYk6hLWon_UP0rMptB41ljfZkhNOvbLUPVxDowH066xA42Zz173_rsIlQAnYpv2qrlmQZ9MjqZVlYcdsWd0ZAN7SEaKwZ40sU9zdkP95PVkKfH4uFw5BkGh4qZanzjr9Y-b7iDM" --header="Connection: keep-alive" "https://doc-14-5k-docs.googleusercontent.com/docs/securesc/lcn000d4f5ncb3531bgn3uus2eb0i5pv/1tucrmhioet355kq68274s05aqbk/1598494200000/03515051603858730688/03515051603858730688/1-SiBIaHKwoSznjRd68TXpjTYyXl52bDM?e=download&authuser=0&nonce=79hmk03b4fg0e&user=03515051603858730688&hash=gf7v8pbslm7srs34pveq58vn05lqhl" -c -O 'tokenization.py'
```

```
--2020-08-27 02:10:38-- https://doc-14-5k-docs.googleusercontent.com/docs/securesc/lcn000d4f5ncb3531bgn3uus2eb0i5pv/1tucrmhioet355kq68274s05aqbk/1598494200000/03515051603858730688/03515051603858730688/1-SiBIaHKwoSznjRd68TXpjTYyXl52bDM?e=download&authuser=0&nonce=79hmk03b4fg0e&user=03515051603858730688&hash=gf7v8pbslm7srs34pveq58vn05lqhl
```

Resolving doc-14-5k-docs.googleusercontent.com (doc-14-5k-docs.googleusercontent.com)...

64.233.188.132, 2404:6800:4008:c06::84

Connecting to doc-14-5k-docs.googleusercontent.com (doc-14-5k-

docs.googleusercontent.com)|64.233.188.132|:443... connected.

HTTP request sent, awaiting response... 200 OK

Length: 17318 (17K) [text/x-python]

Saving to: 'tokenization.py'

tokenization.py 100%[=====>] 16.91K --.-KB/s in 0s

2020-08-27 02:10:38 (86.8 MB/s) - 'tokenization.py' saved [17318/17318]

time: 478 ms

In [ ]:

```
#import tokenization - We have given tokenization.py file
import tokenization as bert_tokenizer
```

time: 13.9 ms

In [ ]:

```
# Create tokenizer " Instantiate FullTokenizer"
# name must be "tokenizer"
# the FullTokenizer takes two parameters 1. vocab_file and 2. do_lower_case
# we have created these in the above cell ex: FullTokenizer(vocab_file, do_lower_case )
# please check the "tokenization.py" file the complete implementation

tokenizer=bert_tokenizer.FullTokenizer(vocab_file, do_lower_case)
```

time: 119 ms

### Grader function 3

In [ ]:

```
#it has to give no error
def grader_tokenize(tokenizer):
    out = False
    try:
        out=('[CLS]' in tokenizer.vocab) and ('[SEP]' in tokenizer.vocab)
    except:
        out = False
    assert(out==True)
    return out
grader_tokenize(tokenizer)
```

Out[ ]:

True

time: 8.9 ms

In [ ]:

```
# Create train and test tokens (X_train_tokens, X_test_tokens) from (X_train, X_test) using Tokenizer and

# add '[CLS]' at start of the Tokens and '[SEP]' at the end of the tokens.

# maximum number of tokens is 55(We already given this to BERT layer above) so shape is (None, 55)

# if it is less than 55, add '[PAD]' token else truncate the tokens length.(similar to padding)

# Based on padding, create the mask for Train and Test ( 1 for real token, 0 for '[PAD]'),
# it will also same shape as input tokens (None, 55) save those in X_train_mask, X_test_mask

# Create a segment input for train and test. We are using only one sentence so all zeros. This shape will also (None, 55)

# type of all the above arrays should be numpy arrays

# after execution of this cell, you have to get
# X_train_tokens, X_train_mask, X_train_segment
# X_test_tokens, X_test_mask, X_test_segment
```

In [ ]:

```
import numpy as np
from keras.preprocessing.sequence import pad_sequences
```

time: 55.2 ms

In [ ]:

```
def tokenization(data):
    X_tokens=np.empty((0,55), int)
    X_mask=np.empty((0,55), int)
    X_segment=np.zeros([len(data),55])
    for idx,ele in enumerate(data):
        token = tokenizer.tokenize(ele)
```



```
time: 21.5 ms
```

time: 11min 54s

```
time: 40.6 s
```

[illegible]

In [ ]:

```
import pickle
```

time: 1.01 ms

In [ ]:

```
##save all your results to disk so that, no need to run all again.
pickle.dump((X_train, X_train_tokens, X_train_mask, X_train_segment,
y_train),open('train_data.pkl','wb'))
pickle.dump((X_test, X_test_tokens, X_test_mask, X_test_segment, y_test),open('test_data.pkl','wb')
)
```

time: 365 ms

In [ ]:

```
#you can load from disk
#X_train, X_train_tokens, X_train_mask, X_train_segment, y_train =
pickle.load(open("train_data.pkl", 'rb'))
#X_test, X_test_tokens, X_test_mask, X_test_segment, y_test = pickle.load(open("test_data.pkl", 'r
b'))
```

time: 458 ms

## Grader function 4

In [ ]:

```
def grader_alltokens_train():
    out = False

    if type(X_train_tokens) == np.ndarray:

        temp_shapes = (X_train_tokens.shape[1]==max_seq_length) and
(X_train_mask.shape[1]==max_seq_length) and \
        (X_train_segment.shape[1]==max_seq_length)

        segment_temp = not np.any(X_train_segment)

        mask_temp = np.sum(X_train_mask==0) == np.sum(X_train_tokens==0)

        no_cls = np.sum(X_train_tokens==tokenizer.vocab['[CLS]'])==X_train_tokens.shape[0]

        no_sep = np.sum(X_train_tokens==tokenizer.vocab['[SEP]'])==X_train_tokens.shape[0]

        out = temp_shapes and segment_temp and mask_temp and no_cls and no_sep

    else:
        print('Type of all above token arrays should be numpy array not list')
        out = False
    assert(out==True)
    return out

grader_alltokens_train()
```

Out[ ]:

True

time: 53.4 ms

## Grader function 5

In [ ]:

```
def grader_alltokens_test():
    out = False
```

```

if type(X_test_tokens) == np.ndarray:

    temp_shapes = (X_test_tokens.shape[1]==max_seq_length) and
(X_test_mask.shape[1]==max_seq_length) and \
    (X_test_segment.shape[1]==max_seq_length)

    segment_temp = not np.any(X_test_segment)

    mask_temp = np.sum(X_test_mask==0) == np.sum(X_test_tokens==0)

    no_cls = np.sum(X_test_tokens==tokenizer.vocab['[CLS]'])==X_test_tokens.shape[0]

    no_sep = np.sum(X_test_tokens==tokenizer.vocab['[SEP]'])==X_test_tokens.shape[0]

    out = temp_shapes and segment_temp and mask_temp and no_cls and no_sep

else:
    print('Type of all above token arrays should be numpy array not list')
    out = False
    assert(out==True)
    return out
grader_alltokens_test()

```

Out[ ]:

True

time: 30.3 ms

## Part-4: Getting Embeddings from BERT Model

We already created the BERT model in the part-2 and input data in the part-3. We will utilize those two and will get the embeddings for each sentence in the Train and test data.

In [ ]:

```
bert_model.input
```

Out[ ]:

```

[<tf.Tensor 'input_word_ids:0' shape=(None, 55) dtype=int32>,
 <tf.Tensor 'input_mask:0' shape=(None, 55) dtype=int32>,
 <tf.Tensor 'segment_ids:0' shape=(None, 55) dtype=int32>]

```

time: 5.12 ms

In [ ]:

```
bert_model.output
```

Out[ ]:

```
<tf.Tensor 'keras_layer/StatefulPartitionedCall:0' shape=(None, 768) dtype=float32>
```

time: 8.52 ms

In [ ]:

```

# get the train output, BERT model will give one output so save in
# X_train_pooled_output
X_train_pooled_output=bert_model.predict([X_train_tokens,X_train_mask,X_train_segment])

```

time: 9min 27s

In [ ]:

```
# get the test output, BERT model will give one output so save in
# X_test_pooled_output
X_test_pooled_output=bert_model.predict([X_test_tokens,X_test_mask,X_test_segment])
```

time: 2min 20s

In [ ]:

```
##save all your results to disk so that, no need to run all again.
pickle.dump((X_train_pooled_output, X_test_pooled_output),open('final_output.pkl','wb'))
```

time: 860 ms

In [ ]:

```
#X_train_pooled_output, X_test_pooled_output= pickle.load(open('final_output.pkl', 'rb'))
```

## Grader function 6

In [ ]:

```
#now we have X_train_pooled_output, y_train
#X_test_pooled_output, y_test

#please use this grader to evaluate
def greader_output():
    assert(X_train_pooled_output.shape[1]==768)
    assert(len(y_train)==len(X_train_pooled_output))
    assert(X_test_pooled_output.shape[1]==768)
    assert(len(y_test)==len(X_test_pooled_output))
    assert(len(y_train.shape)==1)
    assert(len(X_train_pooled_output.shape)==2)
    assert(len(y_test.shape)==1)
    assert(len(X_test_pooled_output.shape)==2)
    return True
grader_output()
```

Out[ ]:

True

time: 13.8 ms

## Part-5: Training a NN with 768 features

Create a NN and train the NN.

1. You have to use AUC as metric.
2. You can use any architecture you want.
3. You have to use tensorboard to log all your metrics and Losses. You have to send those logs.
4. Print the loss and metric at every epoch.
5. You have to submit without overfitting and underfitting.

In [ ]:

```
##imports
from tensorflow.keras.layers import Input, Dense, Activation, Dropout
from tensorflow.keras.models import Model
import keras
from keras import models, layers
from keras.models import Model
from keras.layers import BatchNormalization, Activation, Flatten,Dense,Dropout,Input
from keras.optimizers import Adam,SGD
from keras.callbacks import Callback
from keras.initializers import HeNormal
```

```

from keras.layers import GlobalAveragePooling1D
from keras.layers.convolutional import Conv1D
from keras.layers.convolutional import MaxPooling1D
import tensorflow

```

time: 6.67 ms

In [ ]:

```

input_Layer      = layers.Input(shape=(768,))
dense_layer_1    = layers.Dense(512,activation='relu')(input_Layer)
dropout_Layer    = Dropout(0.2)(dense_layer_1)
dense_layer_2    = layers.Dense(256,activation='relu')(dropout_Layer)
dense_layer_3    = layers.Dense(64,activation='relu')(dense_layer_2)
drop_out_Layer   = Dropout(0.5)(dense_layer_3)
dense_layer_4    = layers.Dense(2,activation='softmax')(drop_out_Layer)
model = Model(inputs=[input_Layer], outputs=[dense_layer_4])

```

time: 84.3 ms

In [ ]:

```
model.summary()
```

Model: "functional\_3"

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[ (None, 768) ]	0
dense (Dense)	(None, 512)	393728
dropout (Dropout)	(None, 512)	0
dense_1 (Dense)	(None, 256)	131328
dense_2 (Dense)	(None, 64)	16448
dropout_1 (Dropout)	(None, 64)	0
dense_3 (Dense)	(None, 2)	130

```

Total params: 541,634
Trainable params: 541,634
Non-trainable params: 0

```

time: 3.89 ms

In [ ]:

```

from google.colab import drive
drive.mount('/content/drive')

```

Go to this URL in a browser: [https://accounts.google.com/o/oauth2/auth?client\\_id=947318989803-6bn6qk8qdgf4n4g3pfee6491hc0brc4i.apps.googleusercontent.com&redirect\\_uri=urn%3aietf%3awg%3aoauth%3a2.0%b&scope=email%20https%3a%2f%2fwww.googleapis.com%2fauth%2fdocs.test%20https%3a%2f%2fwww.googleapis.com%2fauth%2fdrive%20https%3a%2f%2fwww.googleapis.com%2fauth%2fdrive.photos.readonly%20https%3a%2f%2fwww.googleapis.com%2fauth%2fpeopleapi.readonly&response\\_type=code](https://accounts.google.com/o/oauth2/auth?client_id=947318989803-6bn6qk8qdgf4n4g3pfee6491hc0brc4i.apps.googleusercontent.com&redirect_uri=urn%3aietf%3awg%3aoauth%3a2.0%b&scope=email%20https%3a%2f%2fwww.googleapis.com%2fauth%2fdocs.test%20https%3a%2f%2fwww.googleapis.com%2fauth%2fdrive%20https%3a%2f%2fwww.googleapis.com%2fauth%2fdrive.photos.readonly%20https%3a%2f%2fwww.googleapis.com%2fauth%2fpeopleapi.readonly&response_type=code)

Enter your authorization code:

.....

Mounted at /content/drive

time: 22.7 s

In [ ]:

```

import os
os.chdir('/content/drive/My Drive/Applied AI Course/Assignments/27. NLP with Transfer Learning')
!pwd

```

/content/drive/My Drive/Applied AI Course/Assignments/27. NLP with Transfer Learning  
time: 1.01 s

In [ ]:

```
!mkdir -p logs/fit/
```

time: 165 ms

In [ ]:

```
!mkdir save_model
```

time: 160 ms

In [ ]:

```
##create an NN and

def NN(model,X_train_pooled_output,y_train,X_test_pooled_output,y_test):
    auc_metric=tf.keras.metrics.AUC(name="AUC")
    model.compile(loss='categorical_crossentropy',
                  optimizer=Adam(),
                  metrics=['accuracy',auc_metric])
    class Custom_callback(tensorflow.keras.callbacks.Callback):
        def on_epoch_end(self, epoch, logs={}):
            auc = logs.get('AUC')
            if (auc > 0.97) :
                print("\n We have Reached the accepted criteria of AUC -::- {0} , so stopping training!!".
format(auc))
                self.model.stop_training = True

    filepath = "/content/drive/My Drive/Applied AI Course/Assignments/27. NLP with Transfer
Learning/save_model/{epoch:03d}.hdf5"
    model_chkpt = keras.callbacks.ModelCheckpoint(filepath, monitor='val_accuracy', verbose=0, save_b
est_only=True,mode='max')
    log_dir="/content/drive/My Drive/Applied AI Course/Assignments/27. NLP with Transfer
Learning/logs/fit/" + datetime.datetime.now().strftime("%Y%m%d-%H%M%S")
    tensorboard_callback = tf.keras.callbacks.TensorBoard(log_dir=log_dir,histogram_freq=1, write_gra
ph=True,write_grads=True)
    model.fit(X_train_pooled_output, y_train,
              batch_size=64,
              epochs=50,
              verbose=1,
              validation_data=(X_test_pooled_output, y_test),
              callbacks = [model_chkpt,Custom_callback(),tensorboard_callback])
    return model
```

time: 10.6 ms

## Part-6: Creating a Data pipeline for BERT Model

1. Download data from [here](#)
2. Read the csv file
3. Remove all the html tags
4. Now do tokenization [Part 3 as mentioned above]
  - Create tokens,mask array and segment array

5. Get Embeddings from BERT Model [Part 4 as mentioned above] , let it be X\_test

- Print the shape of output(X\_test.shape).You should get (352,768)

6. Predict the output of X\_test with the Neural network model which we trained earlier.

7. Print the occurrences of class labels in the predicted output

</pre>

In [ ]:

```
!wget --header="Host: doc-14-bo-docs.googleusercontent.com" --header="User-Agent: Mozilla/5.0
(Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/84.0.4147.135
Safari/537.36" --header="Accept:
text/html,application/xhtml+xml,application/xml;q=0.9,image/webp,image/apng,*/*;q=0.8,application/
d-exchange;q=0.9" --header="Accept-Language: en-IN,en-GB;q=0.9,en-US;q=0.8,en;q=0.7" --header
="Referer: https://drive.google.com/" --header="Cookie:
AUTH_gnb78hmdiks9t0b8kec09hpa7nncs5e_nonce=8sol7b307osqe; _ga=GA1.2.1804417035.1594643089; NID=20
4=WE4tnDNREeWA-
i7VkeyXZRT0nZVp1JeYYk6hLWon_UP0rMptB4l1jfZkhNOvbLUPVxDowH066xA42Zz173_rsIlQAnYpv2qrlmQZ9MjqZVljYcds
wd0ZAN7SEaKwZ40sU9zdkP95PVkKfH4uFw5BkGh4qZanzjr9Y-b7iDM" --header="Connection: keep-alive"
"https://doc-14-bo-
docs.googleusercontent.com/docs/securesc/lcn000d4f5ncb3531bgn3uus2eb0i5pv/3n701fema2iljh0hh5v4pnun9
0vm/1598496150000/00484516897554883881/03515051603858730688/1QwjqTsqTX2vdy7fTmeXjxP3dq8IAVLpo?e=do
wnload&authuser=0&nonce=8sol7b307osqe&user=03515051603858730688&hash=inqo8dmtikkboj8ic1q8qe4na64fa1
-c -O 'test.csv'
```

```
--2020-08-27 02:43:46-- https://doc-14-bo-
docs.googleusercontent.com/docs/securesc/lcn000d4f5ncb3531bgn3uus2eb0i5pv/3n701fema2iljh0hh5v4pnun9
0vm/1598496150000/00484516897554883881/03515051603858730688/1QwjqTsqTX2vdy7fTmeXjxP3dq8IAVLpo?e=do
wnload&authuser=0&nonce=8sol7b307osqe&user=03515051603858730688&hash=inqo8dmtikkboj8ic1q8qe4na64fa1
```

Resolving doc-14-bo-docs.googleusercontent.com (doc-14-bo-docs.googleusercontent.com)...

64.233.188.132, 2404:6800:4008:c06::84

Connecting to doc-14-bo-docs.googleusercontent.com (doc-14-bo-  
docs.googleusercontent.com)|64.233.188.132|:443... connected.

HTTP request sent, awaiting response... 200 OK

Length: 62100 (61K) [text/csv]

Saving to: 'test.csv'

test.csv 100%[=====>] 60.64K --.-KB/s in 0.005s

2020-08-27 02:43:47 (11.2 MB/s) - 'test.csv' saved [62100/62100]

time: 675 ms

In [ ]:

```
test_csv = pd.read_csv("test.csv")
```

time: 18.3 ms

In [ ]:

```
test_data=test_csv['Text'] # convert dataframe to series pandas
print("Shape of test file:",test_data.shape)
```

Shape of test file: (352,)

time: 1.6 ms

In [ ]:

```
In [ ]:
```

```
# Remove all the html tag
for e in test_data:
    e=re.sub("[<\([\].*?[\]\>]", "", e)
```

time: 3.57 ms

```
In [ ]:
```

```
test_tokens, test_mask, test_segment=tokenization(test_data)
```

time: 240 ms

```
In [ ]:
```

```
test_pooled_output=bert_model.predict([test_tokens,test_mask,test_segment])
```

time: 2.69 s

```
In [ ]:
```

```
test_pooled_output.shape
```

```
Out[ ]:
```

```
(352, 768)
```

time: 3.24 ms

```
In [ ]:
```

```
ytrain_cat = keras.utils.to_categorical(y_train, num_classes=2)
ytest_cat = keras.utils.to_categorical(y_test, num_classes=2)
```

time: 4.68 ms

```
In [ ]:
```

```
%load_ext tensorboard
```

time: 15.3 ms

```
In [ ]:
```

```
import datetime
trained_model=NN(model,X_train_pooled_output,ytrain_cat,X_test_pooled_output,ytest_cat)
```

WARNING:tensorflow:`write\_grads` will be ignored in TensorFlow 2.0 for the `TensorBoard` Callback.

WARNING:tensorflow:`write\_grads` will be ignored in TensorFlow 2.0 for the `TensorBoard` Callback.

Epoch 1/50

1/1250 [.....] - ETA: 0s - loss: 0.5900 - accuracy: 0.7812 - AUC: 0.7679  
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow/python/ops/summary\_ops\_v2.py:1277: stop (from tensorflow.python.eager.profiler) is deprecated and will be removed after 2020-07-01.  
Instructions for updating:  
use `tf.profiler.experimental.stop` instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow/python/ops/summary\_ops\_v2.py:1277: stop (from tensorflow.python.eager.profiler) is deprecated and will be removed after 2020-07-01.  
Instructions for updating:  
use `tf.profiler.experimental.stop` instead.



```
2/1250 [.....] - ETA: 1:14 - loss: 0.5092 - accuracy: 0.8203 - AUC: 0.8275WARNING:tensorflow:Callbacks method `on_train_batch_end` is slow compared to the batch time (batch time: 0.0121s vs `on_train_batch_end` time: 0.1071s). Check your callbacks.
```

```
WARNING:tensorflow:Callbacks method `on_train_batch_end` is slow compared to the batch time (batch time: 0.0121s vs `on_train_batch_end` time: 0.1071s). Check your callbacks.
```

```
1250/1250 [=====] - 8s 6ms/step - loss: 0.3108 - accuracy: 0.8812 - AUC: 0.9400 - val_loss: 0.2566 - val_accuracy: 0.9038 - val_AUC: 0.9652
Epoch 2/50
1250/1250 [=====] - 7s 6ms/step - loss: 0.2642 - accuracy: 0.8958 - AUC: 0.9571 - val_loss: 0.2422 - val_accuracy: 0.8985 - val_AUC: 0.9645
Epoch 3/50
1250/1250 [=====] - 7s 6ms/step - loss: 0.2558 - accuracy: 0.8986 - AUC: 0.9599 - val_loss: 0.2355 - val_accuracy: 0.9020 - val_AUC: 0.9667
Epoch 4/50
1250/1250 [=====] - 7s 6ms/step - loss: 0.2468 - accuracy: 0.9012 - AUC: 0.9627 - val_loss: 0.2223 - val_accuracy: 0.9111 - val_AUC: 0.9698
Epoch 5/50
1250/1250 [=====] - 9s 7ms/step - loss: 0.2417 - accuracy: 0.9027 - AUC: 0.9643 - val_loss: 0.2230 - val_accuracy: 0.9129 - val_AUC: 0.9697
Epoch 6/50
1250/1250 [=====] - 8s 6ms/step - loss: 0.2385 - accuracy: 0.9036 - AUC: 0.9652 - val_loss: 0.2148 - val_accuracy: 0.9140 - val_AUC: 0.9720
Epoch 7/50
1250/1250 [=====] - 7s 6ms/step - loss: 0.2351 - accuracy: 0.9047 - AUC: 0.9663 - val_loss: 0.2326 - val_accuracy: 0.9082 - val_AUC: 0.9701
Epoch 8/50
1250/1250 [=====] - 7s 6ms/step - loss: 0.2314 - accuracy: 0.9074 - AUC: 0.9672 - val_loss: 0.2188 - val_accuracy: 0.9094 - val_AUC: 0.9707
Epoch 9/50
1250/1250 [=====] - 7s 6ms/step - loss: 0.2303 - accuracy: 0.9078 - AUC: 0.9675 - val_loss: 0.2286 - val_accuracy: 0.9050 - val_AUC: 0.9688
Epoch 10/50
1250/1250 [=====] - 7s 6ms/step - loss: 0.2304 - accuracy: 0.9076 - AUC: 0.9675 - val_loss: 0.2285 - val_accuracy: 0.9122 - val_AUC: 0.9691
Epoch 11/50
1250/1250 [=====] - 8s 6ms/step - loss: 0.2286 - accuracy: 0.9079 - AUC: 0.9681 - val_loss: 0.2158 - val_accuracy: 0.9161 - val_AUC: 0.9715
Epoch 12/50
1250/1250 [=====] - 8s 6ms/step - loss: 0.2265 - accuracy: 0.9082 - AUC: 0.9687 - val_loss: 0.2113 - val_accuracy: 0.9166 - val_AUC: 0.9732
Epoch 13/50
1250/1250 [=====] - 7s 6ms/step - loss: 0.2266 - accuracy: 0.9084 - AUC: 0.9685 - val_loss: 0.2122 - val_accuracy: 0.9161 - val_AUC: 0.9726
Epoch 14/50
1250/1250 [=====] - 8s 6ms/step - loss: 0.2244 - accuracy: 0.9100 - AUC: 0.9692 - val_loss: 0.2134 - val_accuracy: 0.9114 - val_AUC: 0.9724
Epoch 15/50
1250/1250 [=====] - 8s 6ms/step - loss: 0.2238 - accuracy: 0.9100 - AUC: 0.9693 - val_loss: 0.2081 - val_accuracy: 0.9179 - val_AUC: 0.9736
Epoch 16/50
1250/1250 [=====] - 8s 6ms/step - loss: 0.2217 - accuracy: 0.9104 - AUC: 0.9699 - val_loss: 0.2078 - val_accuracy: 0.9165 - val_AUC: 0.9736
Epoch 17/50
1250/1250 [=====] - 7s 6ms/step - loss: 0.2236 - accuracy: 0.9093 - AUC: 0.9694 - val_loss: 0.2045 - val_accuracy: 0.9154 - val_AUC: 0.9744
Epoch 18/50
1241/1250 [=====>.] - ETA: 0s - loss: 0.2193 - accuracy: 0.9112 - AUC: 0.9706
We have Reached the accepted criteria of AUC -::- 0.9705479741096497 , so stopping training!!
1250/1250 [=====] - 7s 6ms/step - loss: 0.2194 - accuracy: 0.9112 - AUC: 0.9705 - val_loss: 0.2100 - val_accuracy: 0.9150 - val_AUC: 0.9741
time: 2min 18s
```

In [ ]:

```
%tensorboard --logdir logs/fit
```

Output hidden; open in <https://colab.research.google.com> to view.

In [ ]:

```
y_test_csv=trained_model.predict(test_pooled_output,verbose=1)
```

11/11 [=====] - 0s 2ms/step  
time: 115 ms

In [ ]:

```
y_classes = np.argmax(y_test_csv,axis=1)
```

time: 1.03 ms

In [ ]:

```
counts=np.unique(y_classes,return_counts=True)  ## Returns sorted unique values in the form of Tuple.
```

time: 2.01 ms

In [ ]:

```
print("Zero (",counts[0][0],") Count:",counts[1][0])  
print("One (",counts[0][1],") Count:",counts[1][1])
```

Zero ( 0 ) Count: 26  
One ( 1 ) Count: 326  
time: 4.12 ms

## Summary

In [ ]:

```
from prettytable import PrettyTable  
x = PrettyTable()  
x.field_names = ["Training AUC","Test / Validation AUC"]  
x.add_row([0.9705,0.9741 ])  
print(x)
```

```
+-----+-----+  
| Training AUC | Test / Validation AUC |  
+-----+-----+  
|    0.9705    |         0.9741         |  
+-----+-----+  
time: 13.8 ms
```

## Observations and Steps Followed

### Preprocessing

1. User tensorflow 2.0 to train the model.
2. Downloaded and extracted the Reviews.csv from “Amazon Fine Food Reviews” Kaggle website.
3. We see there are 9 columns in total but as per instructions we took only two columns[‘Text’,‘Score’] for assignment.
4. Dropped the null rows and scored == 3
5. Manipulated the score columns into 1 and 0 from range(0,5)
6. updated the dataset by removing all the rows where the length of words in text is greater than 50.
7. Removed all HTML tags from text column.
8. split the data into train and test with 80% train.
9. As the count plots work on Data frame, converted the targets into data frame and plotted the counts in bar graph.

Source :: <https://medium.com/@aieeshashafique/feature-extraction-from-bert-25887ed2152a>

**Creating a BERT model from the Tensorflow HUB.**

1. We can choose any max length, BERT has a constraint on the maximum length of a sequence after tokenizing. For any BERT model, the maximum sequence length after tokenization is 512. But we can set any sequence length equal to or below this value.
2. Created three input layers with the names [input\_word\_ids , input\_mask, segment\_ids]
3.
  - input token ids (tokenizer converts tokens using vocab file)
  - input masks (1 for useful tokens, 0 for padding)
  - segment ids (for 2 text training: 0 for the first one, 1 for the second one)
1. Created a model with concatenation of three input layers and pooled output layer (pooled output of shape [batch\_size, 768] with representations for the entire input sequences)

### 3. Tokenization

1. Install "sentencepiece" library to import tokenization.py file.
2. Import tokenizer using the original vocab file, do lower case all the word pieces and then tokenize the sentences.
3. Defined function "tokenization" to generate Id's, segments and masks based on the original BERT
4. Creating tokens, mask, and segment by using tokenization function for X\_train and X-test.

### 4. getting the pretrained embedding Vector for a given review from the BERT.

1. Generating the X\_train\_pooled\_output and X\_test\_pooled\_output Embeddings using the pre-trained model using Id's, segments and masks.

### 5. Using the embedding data apply NN and classify the reviews.

1. As the "test\_pooled\_output" shape is (352, 768) , here we are defining the model with input layer shape of (768,) followed with dense layers and dropout's.
2. Creating a neural network function with 5 parameters NN(model,X\_train\_pooled\_output,y\_train,X\_test\_pooled\_output,y\_test)
3. Defining the callbacks based on few conditions.

### 6. Creating a Data pipeline for BERT Model.

1. Downloading test data
2. Creating the test Id's , mask and segments
3. Converting into categories.
4. Training the model with train pooled data and Y\_categories.
5. Predicting the values for y\_test
6. defining the y\_classes for the y\_test with numpy library argmax function.
7. Counting the predicted values (counts=np.unique(y\_classes,return\_counts=True) ## Returns sorted unique values in the form of Tuple. )